

# Predicting a vehicles velocity using dashcam footage

## A deep learning approach

Florian Wolf, Department of Mathematics and Statistics  
Franz Herbst, Department of Physics

Machine Learning using Matlab  
Universität Konstanz

January 19, 2021

# Table of content

- 1 Motivation, data collection and initial assumptions
- 2 Preprocessing and optical flow
- 3 Method selection and architecture
- 4 Fine-tuning of the model
  - Initial tuning
  - Problems and possible solutions
  - Simplified model
  - Siamese approach: flow field and frame (new splitting)
- 5 Current and further work

# The “comma ai speed challenge”<sup>1</sup>

## Motivation

- HERE ARE SOME MOTIVATIONAL WORDS NEEDED

## Data collection:

- “comma ai speed challenge” provides two videos:
  - Train video: 24000 frames, shoot at 20 frames per second, including ground truths
  - Test video: 10798 frames, shoot at 20 frames per second, no ground truths, used to applications
- Split train video after 80% with hard cut off (ability the generalize), to get train and test datasets

## Initial assumptions

- Use mean squared error (MSE) as a performance measure
- How to evaluate a prediction? Assumptions:
  - $MSE \leq 10$ : good
  - $MSE \leq 5$ : better
  - $MSE \leq 3$ : correct

---

<sup>1</sup><https://github.com/commaai/speedchallenge>

# Preprocessing

- Frame size of  $(640, 480, 3)$  pixels
- Cut off last 60 pixels, to remove black frame inside the car
- Sample down the frame to half its size, due to computational limitations



Original frame



Cut off the last 60 pixels, downsampled

# Optical flow using “Farneback pyramid method” [2]

- Global method to solve the optical flow equation

$$\partial_x f \cdot V_x + \partial_y f \cdot V_y + \partial_t f = 0$$

for an image sequence  $(f_t)_t$  with  $f_t : \Omega \rightarrow \mathbb{R}^3$ , for all  $t$ , and the (dense) flow field  $V : \Omega \rightarrow \mathbb{R}^2, \omega \mapsto (V_x(\omega), V_y(\omega))$ .

- Uses a downsampling pyramid, to solve the equation for different resolutions of the image
- Parameters for the Farneback method

pyramid levels := 3

pyramid scaling := 0.5

window size := 6

pixel neighborhood size := 5

SD of the gaussian filter := 1.1

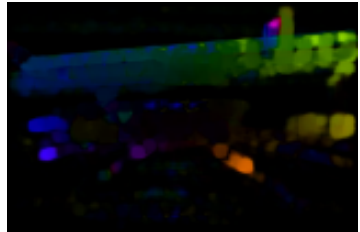
- Result: **Flow field with (160, 105, 3) pixels**

# Visualization of the flow field

- Flow field is a two-dimensional vector field
- RGB representation via
  - Transform flow field into polar coordinates  $(V_x, V_y) \mapsto (r, \varphi)$
  - Normalize magnitudes  $r$  for the third channel
  - Values of the second channel are all set to 255
  - Multiply angle  $\varphi$  by factor  $\frac{180}{2\pi}$  for the first channel
- Sample down the resolution again, to speed up the training



Original frame



Corresponding flow field, already sampled down

# Convolutional neural network and initial architecture

## Method selection

- Speed prediction is a **non-linear regression** task  $\rightsquigarrow$  Neural network
- Use convolution layers to perform feature extraction  $\rightsquigarrow$  **convolutional neural network** (CNN)

## Initial architecture

- Paper of NVIDIA work group [1] of a CNN for self-driving cars
- Enough complexity and layers to handle the task and lots of possibilities to fine-tune it
- Initial results with the raw model: MSE of under 3 on the training set and around 18-20 on the testing set  
 $\Rightarrow$  Improvements needed

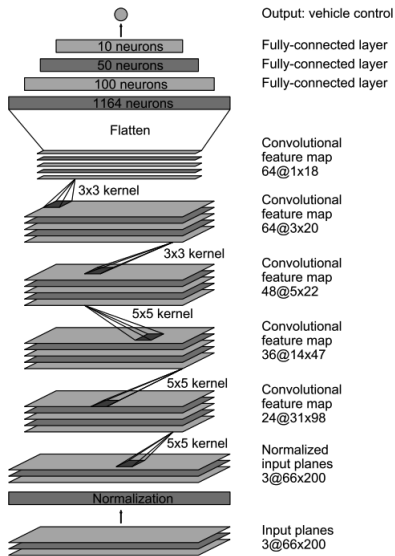


Figure: Original architecture of the NVIDIA paper [1].



# Batch Normalization, Dropout layers, activation function and pooling

- Batch normalization to speed up the training [3]
- Initial activation function:  $\text{ReLU} : \mathbb{R} \rightarrow \mathbb{R}_0^+, x \mapsto \max\{0, x\}$ , still MSE of over 15 on the testing set, less than 2 on the training set  
 $\Rightarrow$  Overfitting problems
- Found paper about dropout layers [4] to reduce overfit
- Solve problems of dead neurons using

$$\text{leakyReLU} : \mathbb{R} \rightarrow \mathbb{R}, x \mapsto \begin{cases} x, & x \geq 0 \\ c \cdot x, & x < 0 \end{cases}$$


with  $c = 0.01$ , MSE of around 11 on the testing set and less than 3 on the training set

# Problems

We identified three possible problems for our poor results

- 1 Too complex model, as initially used for autonomous driving or insufficient amount of information put into the model
- 2 Brightnesses/illumination changes in the frames, therefore unstable calculations of the optical flow
- 3 Too ambiguous splitting, as the training and testing datasets represent totally different road traffic scenarios

# New splitting: situational splitting



**Figure:** Speed distribution in the training video, including labels for different scenarios.

# Possible solutions

- 1 **Simplify model:** pooling layers (maximum and average pooling) to get more compression  
**Siamese approach:** put flow field and raw frame into the model or put two consecutive frames into the model
- 2 **Add additional noise:** add noise before computing the optical flow field, to get more invariance regarding illumination changes
- 3 **Different splitting:** get better ratio between different scenarios, by using different data splittings: finer one and a more specific one based on the different road traffic situations in the video

# Pooling layers (initial splitting)

Initial splitting, 8 epochs	ReLU		leakyReLU	
	Train	Test	Train	Test
No pooling	2.85	12.08	2.45	10.75
Max pooling	5.62	11.82	5.52	10.29
Max pooling (15 epochs)	-	-	<b>3.22</b>	<b>9.63</b>
Average pooling	7.70	11.40	6.08	13.09

**Table:** MSE results of the network using different pooling strategies, one dropout layers, two different activation functions and the initial splitting. We trained each of the models for eight epochs.

# Siamese approach (new splitting)

RESULTS ARE NEEDED :)

# Current work

- 1 Add additional noise to frames
- 2 Siamese approach for two consecutive frames
- 3 Try to figure out why the new splitting does not work

# Further ideas: Continuous modelling assumption

THIS IS JUST A TEST



# References



Mariusz Bojarski et al. “End to End Learning for Self-Driving Cars”. In: (Apr. 2016). URL: <https://arxiv.org/pdf/1604.07316v1.pdf>.



Gunnar Farnebäck. “Two-Frame Motion Estimation Based on Polynomial Expansion”. In: *Scandinavian Conference on Image Analysis* (2003), pp. 363–370.



Sergey Ioffe and Christian Szegedy. “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift”. In: (Feb. 2015). URL: <https://arxiv.org/pdf/1502.03167.pdf>.



Nitish Srivastava et al. “Dropout: A Simple Way to Prevent Neural Networks from Overfitting”. In: *Journal of Machine Learning Research* 15.56 (2014), pp. 1929–1958. URL: <http://jmlr.org/papers/v15/srivastava14a.html>.