

Predicting a vehicles velocity using dashcam footage

A deep learning approach

Florian Wolf, Department of Mathematics and Statistics
Franz Herbst, Department of Physics

Machine Learning using Matlab Python
Universität Konstanz

February 1, 2021

Table of content

- 1 Motivation and initial Dataset**
- 2 Analysis of the dataset**
- 3 Preprocessing using optical flow**
- 4 Method selection and architecture**
- 5 Different optimization strategies**
 - Change components of the initial architecture
 - Siamese approach
 - Problems and possible solutions
- 6 Current and further work**
 - Acquire more data
 - Additional noise
- 7 Summary and application ideas**

The “comma ai speed challenge”¹

Motivation

- autonomous driving is currently one of the most prominent problems in machine learning
- but quite hard to set up on a desktop pc
- predicting a vehicles velocity from video footage is a related but also much more simplified task

Initial Dataset:

- training video with 20400 frames (20 fps)
- data file with velocity of the car at each frame
- test video with 10798 frames (20 fps)

Evaluation:

- the mean squared error (MSE) is used to measure performance

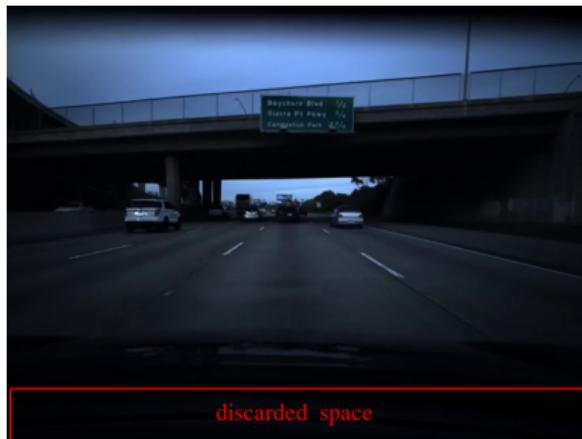
$$\mathcal{L} = \sum_i (p(x_i) - y_i)^2$$

¹<https://github.com/commaai/speedchallenge>

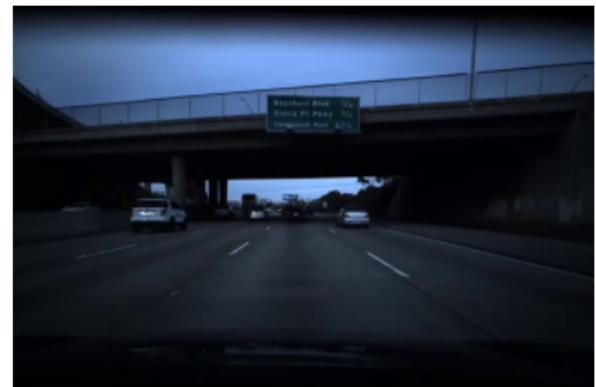
Analysis of the dataset

Video data:

- frame size of (640, 480, 3) pixels
- cut off last 60 pixels, to remove black frame inside the car
- sample down the frame to half its size, to reduce computation time



Original frame



Cut off the last 60 pixels, downsampled

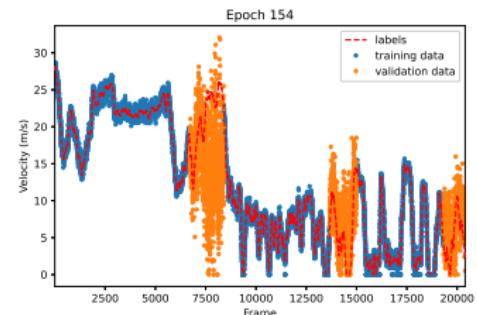
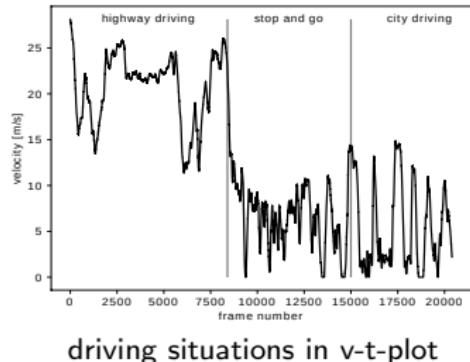
Analysis of the dataset

Splitting of the dataset

- Initial splitting: hard cut off after 80% of the frames
- Situational splitting: divide dataset into blocks of different driving scenarios, splitting with 80% test and 20% validation data on each

Evaluation:

- variance $\sqrt{\mathcal{L}} \gtrsim 16$: no fitting
- $10 \lesssim \sqrt{\mathcal{L}} \lesssim 16$: average velocity fitted
- $5 \lesssim \sqrt{\mathcal{L}} \lesssim 10$: qualitative detection
- $1 \lesssim \sqrt{\mathcal{L}} \lesssim 5$: quantitative detection
- $\sqrt{\mathcal{L}} \lesssim 1$: perfect detection



example performance on training set
training: $\sqrt{\mathcal{L}} = 0.4$, test: $\sqrt{\mathcal{L}} = 6.3$

Optical flow using “Farneback pyramid method” [2]

- Global method to solve the optical flow equation

$$\partial_x f \cdot V_x + \partial_y f \cdot V_y + \partial_t f = 0$$

for an image sequence $(f_t)_t$ with $f_t : \Omega \rightarrow \mathbb{R}^3$, for all t , and the (dense) flow field $V : \Omega \rightarrow \mathbb{R}^2$, $\omega \mapsto (V_x(\omega), V_y(\omega))$.

- Uses a downsampling pyramid, to solve the equation for different resolutions of the image
- Parameters for the Farneback method

pyramid levels := 3

pyramid scaling := 0.5

window size := 6

pixel neighborhood size := 5

SD of the gaussian filter := 1.1

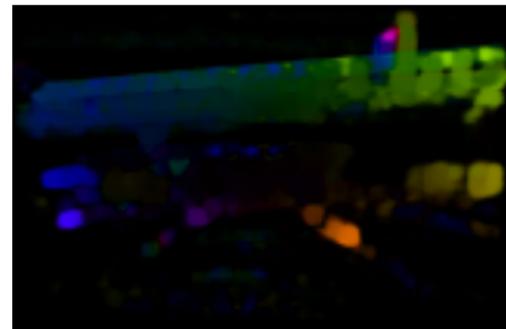
- Result: **Flow field with (160, 105, 3) pixels**

Visualization of the flow field

- Flow field is a two-dimensional vector field
- RGB representation via
 - Transform flow field into polar coordinates $(V_x, V_y) \xrightarrow{\sim} (r, \varphi)$
 - Normalize magnitudes r for the third channel
 - Values of the second channel are all set to 255
 - Multiply angle φ by factor $\frac{180}{2\pi}$ for the first channel
- Sample down the resolution again to speed up the training



Input frame



Corresponding flow field

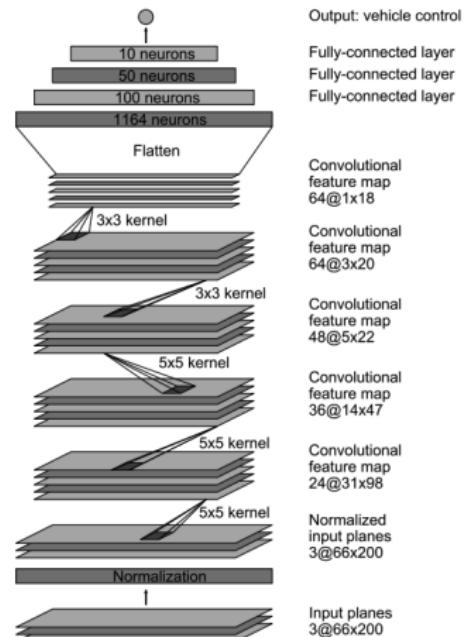
Choosing an initial architecture

Method selection

- speed prediction is a non-linear regression task \Rightarrow neural network
- task involves feature extraction \Rightarrow convolutional neural network (CNN)

Initial architecture

- using paper of *NVIDIA* work group [1] of a CNN for self-driving cars adapted on our initial data
- enough complexity and layers to handle the task and lots of possibilities to fine-tune it
- Initial results with the raw model: $\mathcal{L} < 3$ on the training set and about $\mathcal{L} \approx 19$ (using an old splitting) on the test set



Original architecture of the *NVIDIA* paper [1]

Our approaches to optimize the results

1 change components of the initial architecture

- adding different pooling layers
- use other activation functions

2 change the architecture

- expand structure to Siamese network
- use different setups

3 change the input data

- use brightness augmentation
- acquire more data

Batch Normalization, Dropout layers, activation function and pooling

- Batch normalization to speed up the training [3]
- Initial activation function: $\text{ReLU} : \mathbb{R} \rightarrow \mathbb{R}_0^+, x \mapsto \max\{0, x\}$, still MSE of over 15 on the testing set, less than 2 on the training set
⇒ Overfitting problems
- Dropout layers [4] to make the model more robust and reduce overfitting
- Solve problems of dead neurons using

$$\text{leakyReLU} : \mathbb{R} \rightarrow \mathbb{R}, x \mapsto \begin{cases} x, & x \geq 0 \\ c \cdot x, & x < 0 \end{cases}$$

with $c = 0.01$, MSE of around 11 on the testing set and less than 3 on the training set

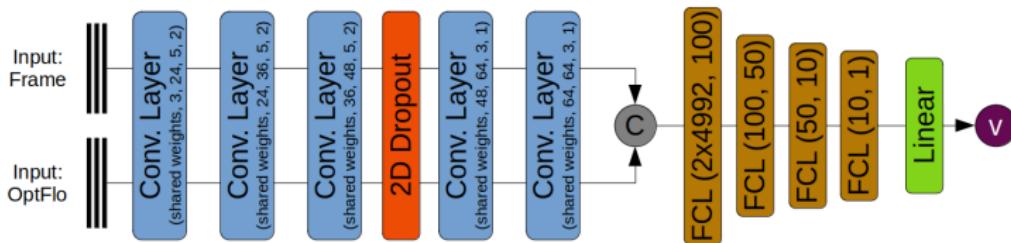
Pooling layers (initial splitting)

Initial splitting, 8 epochs	ReLU		leakyReLU	
	Train	Test	Train	Test
No pooling	2.85	12.08	2.45	10.75
Max pooling	5.62	11.82	5.52	10.29
Max pooling (15 epochs)	-	-	3.22	9.63
Average pooling	7.70	11.40	6.08	13.09

Table: MSE results of the network using different pooling strategies, one dropout layer, two different activation functions and the initial splitting. We trained each of the models for eight epochs.

Siamese Architecture: Setup

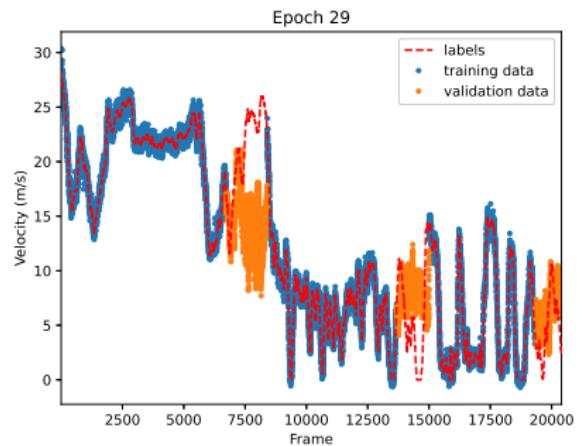
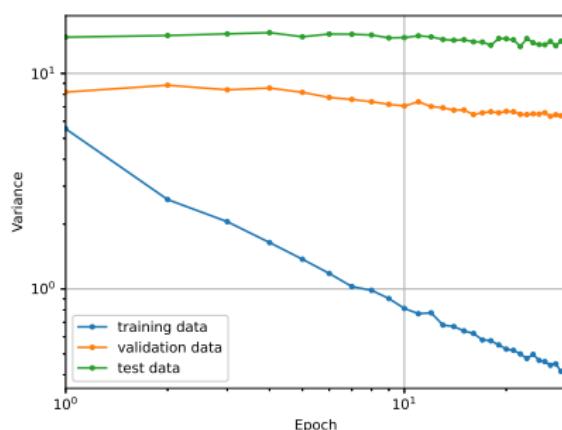
- based on the initial architecture
- using the same convolutional layers on raw frames and optical flow
- weighted sum of the results into fully connected layers



- Use model also for two consecutive frames f_t and f_{t+1}
⇒ advantage: no previous calculation of optical flow needed

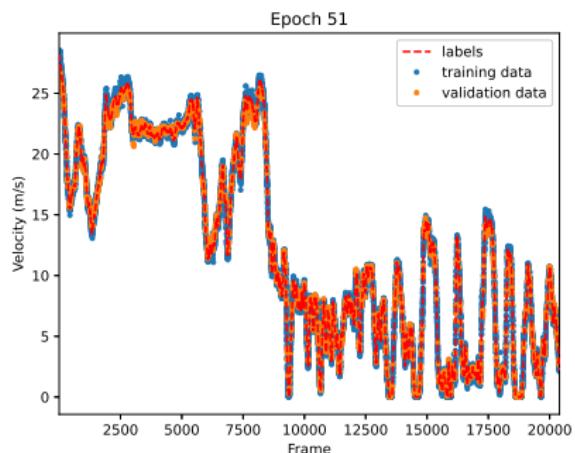
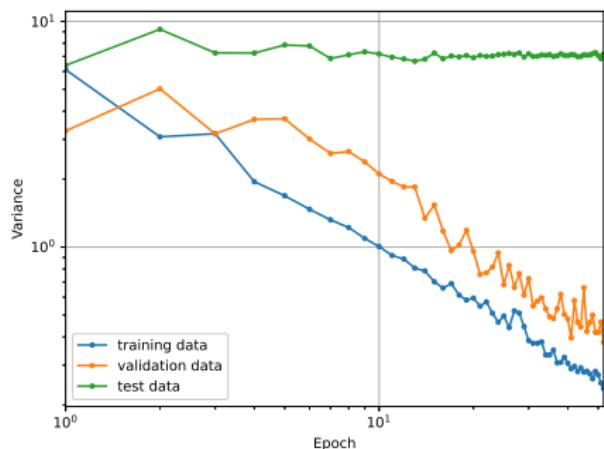
Siamese Architecture: Performance

Performance (Siamese network frame with optical flow):



Siamese Architecture: Performance

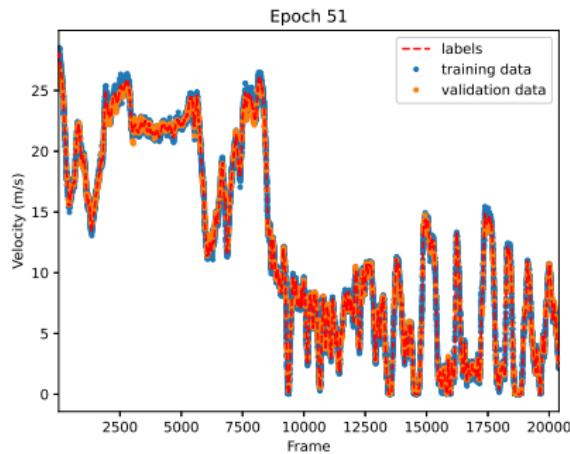
Performance (Siamese network with two frames):



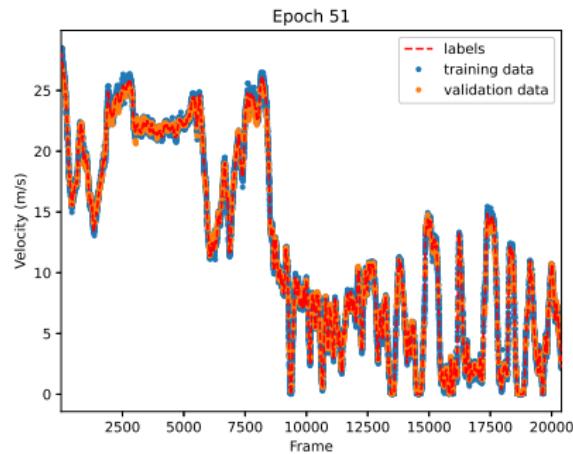
Problems

We identified two problems with our approaches

- 1 overfitting, when validation data is not randomly distributed in the data but a of continuous blocks
- 2 predictions are very susceptible for brightnesses/illumination changes in the frames, because of unstable calculations of the optical flow



random shuffled validation data



block validation data

Problems and possible solutions

Problems

- Predictions are limited to data very similar to the training data
- Predictions do not generalize well

Most likely explanations:

- Very limited dataset for a complex model
- Network is not prepared for unseen situations
- Changes in brightness and illumination

Solution:

- Acquire more data for various driving situations and/or use a less complex model
- Add noise to the data, more robustness in the optical flow calculation

Acquire more data and evaluation of the test video

Method

- video producing and velocity detection with common apps
 - *open street maps*: GPS tracking (.gpx-files)
 - *open camera*: dashcam footage

Advantages and Issues

- very easy method to create a lot of data (if a car is available)
- velocity has some uncertainties (needs to be extrapolated to cover all frames; uncertainties of GPS vs. car sensors)
- frame rate differs a little from original data (24 fps vs 20 fps)

Contrast and brightness augmentation

- Additional noise to frames **before** calculating the flow field.
- Change the brightness and contrast of an image via

$$\text{frame}_{\text{augmented}}(i, j) = \alpha(i, j) \cdot \text{frame}(i, j) + \beta(i, j)$$

with functions α (contrast: > 1 increase, < 1 decrease) and β (brightness).
To get some noise into the frames, we used

$$\begin{aligned}\alpha &\sim \mathcal{U}(0, 1) + 0.35 \\ \beta &\sim \mathcal{U}(-5, 35),\end{aligned}$$

where $\mathcal{U}(a, b)$ is the uniform distribution in an interval $[a, b]$ for $a < b$.

Usability in real world application

- using a similar setup as for acquiring the test data it would be possible to achieve a live prediction
- to test our predictions we wrote a method reading a video and predicting the speed
- even with calculating the optical flow we reached a speed of 30 fps \Rightarrow live prediction would be possible
- at the moment there is no completely convincing model
- a lot of programming would be involved

References

- [1] Mariusz Bojarski et al. "End to End Learning for Self-Driving Cars". In: (Apr. 2016). URL: <https://arxiv.org/pdf/1604.07316v1.pdf>.
- [2] Gunnar Farnebäck. "Two-Frame Motion Estimation Based on Polynomial Expansion". In: Scandinavian Conference on Image Analysis (2003), pp. 363–370.
- [3] Sergey Ioffe and Christian Szegedy. "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift". In: (Feb. 2015). URL: <https://arxiv.org/pdf/1502.03167.pdf>.
- [4] Nitish Srivastava et al. "Dropout: A Simple Way to Prevent Neural Networks from Overfitting". In: Journal of Machine Learning Research 15.56 (2014), pp. 1929–1958. URL: <http://jmlr.org/papers/v15/srivastava14a.html>.