A machine learning approach to predict a vehicles velocity using dashcam video footage

Florian Wolf

Department of Mathematics and Statistics, University Konstanz
Konstanz, Germany
florian.2.wolf@uni-konstanz.de

Franz Herbst

Department of Physics, University Konstanz
Konstanz, Germany
franz.herbst@uni-konstanz.de

Abstract—**Calculating the velocity of a moving camera relative to its surrounding, the so-called visual odometry problem, is a very challenging task and heavily studied in the area of computer vision. Especially in the field of self-driving cars, a fast and dependable velocity calculation is a high priority. In this report we will give a machine learning approach to solve the visual odometry problem, using optical flow fields combined with convolutional neuronal networks, as well as siamese neuronal networks. We use a data set with real world dashcam footage and even extend the data by gathering our own driving videos.**

Index Terms—**deep learning, computer vision, visual odometry, convolutional neuronal network, siamese network, optical flow**

I. INTRODUCTION

Autonomous driving is considered to have a key role in the future of human mobility. Recently the topic has therefore gained great attention in research, economy and politics [1]. In this project we approached this topic in a related but more simplified setup and tried out several machine learning techniques and network architectures.

A. Aim of the project

In this project we tried to predict the velocity of a car based on a dashcam video. This was inspired by the *comma*[1] speed challenge[2] published in 2018. To achieve this goal we used optical flow analysis to preprocess the data and neuronal networks with different structures (classical and siamese) to predict the velocity. In order to measure the quality of the predictions we used their mean squared errors of the measured velocity.

II. DATA COLLECTION, ANALYSIS AND PREPROCESSING

To start the project we used the database given in the *comma* speed challenge project, which consists of a 17 minute training video (20400 frames) and the corresponding car velocity, as well as a 9 minute test video without labels to validate the model on an unknown data set.

A. Data collection

Due to our limited computational power, we mainly used this data set to study different training techniques and model architectures. Still we could not expect good generalization

for this rather small pool of training data. We therefore developed a technique to acquire more data that needed minimal resources. We used the open source smartphone apps *open camera* and *open street maps* to cast a video while driving and map the velocity using GPS at the same time. So we were able to train our final model with 1 hour 40 minutes of video data.

B. Data analysis

In order to evaluate our raw data with the model, we analysed our dataset concerning the recorded images and driving scenarios. The original frames have a size of $640 \times 380 \times 3$ pixels (RGB). To reduce the computation time, we cut the borders of the frames to exclude parts not needed for detection and sampled them down to half of their pixel size.

To visualize the training process in different driving scenarios, we plotted the velocity in dependence of the frame number (Fig. 1). After additional evaluation of the video we were able to classify three different driving scenarios: *highway driving* with high and relatively steady velocities, *stop and go* with low and fluctuating velocities as well as *city driving* with very abrupt speed changes between medium and very low velocities. Every driving scenario is represented with about one third of the training data.


Fig. 1: Distribution of the velocities of all frames in the training video, including the three categories.

We would then evaluate the predictions of each model $p(x_i)$ to the input data x_i using the mean squared error

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^{N} (p(x_i) - y_i)^2 \qquad (1)$$

to real value of the velocity y_i. To classify the resulting value, we used the velocity plot and the following rough classification[3]:

$$
\begin{aligned}
\sqrt{\mathcal{L}} &\gtrsim 16: && \text{no fitting} \\
16 \gtrsim \sqrt{\mathcal{L}} &\gtrsim 10: && \text{average velocity fitted} \\
10 \gtrsim \sqrt{\mathcal{L}} &\gtrsim 5: && \text{scenarios detected} \\
5 \gtrsim \sqrt{\mathcal{L}} &\gtrsim 1: && \text{variances within scenarios detected} \\
1 \gtrsim \sqrt{\mathcal{L}}: && && \text{perfect fitting}
\end{aligned}
$$

TABLE I: Classification of the result based on the square root of the MSE.

For the training process we used a splitting of 80% training data and 20% validation data. Initially we did a hard splitting of the entire data, but this neglects the distribution of different driving scenarios in the video. Therefore we went on first splitting the data into the driving scenarios with then assigning training data and validation data on each.

C. Preprocessing

After the previously explained preparation steps, we needed to find a method to extract information about the motion of the vehicle to be able to predict its velocity. There are two ways of doing this: either fit two following frames into a network, which will be discussed later, or calculate the relative motion, the so-called *optical flow*, between two frames and fit the results into a model.

We used the "Farneback pyramid method" [2] to calculate the optical flow. The idea of this method can be split up into two main steps: First, use a polynomial expansion of the image f and second, solve the optical flow equation

$$\partial_x f \cdot V_x + \partial_y \cdot V_y + \partial_t f = 0$$

for different resolutions of the image (pyramid structure). This methods yields a dense optical flow field V with $\operatorname{im}(V) \subset \mathbb{R}^2$.

We used the following parameters for our calculations:

$$
\begin{aligned}
\text{pyramid levels} &:= 3 \\
\text{pyramid scaling} &:= 0.5 \\
\text{window size} &:= 6 \\
\text{pixel neighborhood size} &:= 5 \\
\text{SD of the gaussian filter} &:= 1.1.
\end{aligned}
$$

We choose three pyramid levels, because we wanted the calculations to be more accurate. To decrease the training duration, we halved the size of the optical flow frames again, resulting in a resolution of $(160, 105, 3)$ pixels per frame. As

[3]These rules do only apply on datasets with an equal distribution of all driving scenarios

we used a window size of six pixels, a comparison between the original optical flow and the down sampled one lead to the result, that we do not loose a lot of information.

The optical flow calculation returns the magnitude and the angle of the flow vectors (overall a two-dimensional vector field), which we transformed into polar coordinates.

To get an RGB image representing the optical flow of two consecutive frames, we normalized the magnitudes and put them into the third channel of the frame. The values of the second channel were all set to the value 255. We then multiplied the angle with the factor $180/(2\pi)$ and set this value for the first channel.

We wanted to see, if the model performs better using the dashcam frames as additional material. Therefore, we did the same down sampling with the raw frames.

III. METHOD SELECTION AND ARCHITECTURE

The prediction of the vehicles speed is a non-linear regression task, so the choice of a neural network is reasonable. Recent architectures (ResNet, GoogLeNet, etc.) have shown that using multiple stacked convolution layers combined with stacked dense layers, perform well on image classification tasks. Therefore the choice of a convolutional neural network is justified.

A. Initial Network

As an initial architecture we used the model presented by the *NVIDIA* work group in [3]. This network was designed for self-driving cars, so the model has enough complexity to handle a task like ours and allowed various possibilities to fine-tune and improve the architecture. The raw structure is shown in Fig. 2.


Fig. 2: Initial Network using a series of convolutional layers. The resulting feature vector is then flattened and mapped into several fully connected layers.

Using the initial model on the training data with a hard splitting, we achieved a MSE between 18 and 20 in the validation set, while having a MSE of less than 3 on the training set. The initial network therefore has some clear overfitting problems, which we tried to address with our fine tuning.

B. Siamese Network

In order to improve our results we also tried to expanded our model and use the original image as well. The corresponding model is shown in Fig. 3 and inspired by the architecture used in [4].

We could use this architecture in two different ways: feeding two consecutive frames into the model or one image and one

Fig. 3: Siamese architecture based on the fined tuned original network. The two inputs pass several convolutional layers with shared weights and a drop out layer. The resulting feature vectors are concatenated and then mapped into several fully connected layers.

Initial splitting, 8 epochs	ReLU		leakyReLU	
	Train	Test	Train	Test
No pooling	2.85	12.08	2.45	10.75
Max pooling	5.62	11.82	5.52	10.29
Max pooling (15 epochs)	-	-	**3.22**	**9.63**
Average pooling	7.70	11.40	6.08	13.09

TABLE II: MSE results of the network using different pooling strategies, one dropout layers, two different activation functions and the initial splitting. We trained each of the models for eight epochs.

optical flow field. In order to keep the structure correspondent, we sampled down the regarding images to the size of the optical flow field.

IV. FINE TUNING

A. Batch normalisation and activation functions

Similar to the lecture, we included batch normalization layers [5], to speed up the training and improve the networks performance.

We also tested different activation functions. As proposed in the lecture, we initially used the ReLu function

$$\text{ReLu} : \mathbb{R} \to \mathbb{R}_0^+, x \mapsto \max\{0, x\}.$$

Using the ReLU function and 15 epochs for training, we achieved a MSE of around 15 on the testing set. We ran the code multiple times, to ensure this result holds. This result was not really promising, so we wanted to decrease the error by modifying the model even more.

To solve the problem of dead neurons[4] of the ReLU function, we tried the leakyReLU function

$$\text{leakyReLU} : \mathbb{R} \to \mathbb{R}, x \mapsto \begin{cases} x, x \geq 0 \\ c \cdot x, x < 0 \end{cases}$$

with a hyperparameter $c = 0.01$. Using the leakyReLU function, we achieved a MSE of around 12 on the testing set and under 3 on the training set. All of the models were trained with only eight epochs, as we still had problems with overfitting the data.

B. Dropout Layer

As we still had a lot of overfitting issues, we decided to include a dropout layer, according to [6]. We tried different positions and different amounts of dropout layers, but using one layer with a probability of $p = 0.5$ after the third convolutional layer seemed to work best.

[4]One can clearly see in the definition of the ReLU function, that neurons with a value below zero cannot participate in the learning process.

C. Pooling layers with initial splitting

We added two generic pooling layers to the network, to reduce the number of parameters[5] of the model and therefore force the network to compress the data even further. One after the second convolutional layer and the second one right before the fully connected layers start. We tested maximum and average pooling with the following parameters

$$\text{kernel size} := 2 \times 2$$
$$\text{stride} := 2$$
$$\text{padding} := 1$$
$$\text{dilatation} := \text{None}$$

The implementation of pooling layers helped a lot, as now the loss on the train and test data seemed to decrease nearly equally. Our results with the initial splitting are shown Table II. We gave the network with max pooling a try with 15 epochs, as the loss on the train and test set was decreasing in a pretty stable manner. Using this network, we achieved for the first time a MSE of under 10 on the test set, which is, according to our assumptions in Section I-A, a pretty good result — especially, as we trained the model mostly on highway scenes and tested it only in city driving scenarios.

D. Siamese approach

As already mentioned in Section III-B, we used the optimized initial network and extended it to a siamese network. Therefore we used shared weights for the convolutional layers, concatenated the two resulting feature vectors and mapped them into the fully connected layers. The results are shown in REFERENCE MISSING.

E. Optimizer and scheduler

For all of our networks we used the *ADAM* [7] algorithm as a (stochastic) optimizer. For scheduling, we tried the *ReduceLROnPlateau*-scheduler of the *PyTorch*[6] module and later switched to the *MultiStepLR*-scheduler.
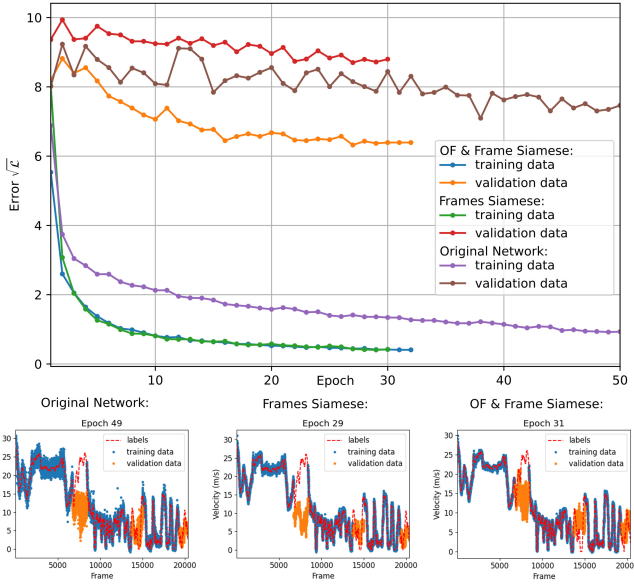
Fig. 4: Comparison of the results using the models: tuned original network and siamese model with frames and optical flow; *top*: training and validation error per epoch; *bottom*: predictions after the last epoch with correct values in red

V. ERROR ANALYSIS AND RESULTS

We identified three possible reasons for our poor results:

(i) predictions are limited to specific video parameters: predictions are bound to camera perspective, changes in brightness
(ii) lack of generalization: perfect fitting on randomly shuffled validation data, qualitative fitting for block splitting and rough fitting for different videos
(iii) very complex model trained with very little data

We came up with the following approaches to solve these problems

(i) We wanted to try adding some additional noise into the frames before calculating the optical flow, to make the calculation more robust against brightness changes. As intentionally adding noise to a frame is quite atypical in computer vision, this idea looked quite interesting.
(ii) Use another splitting. To get a better ratio between highway and city driving scenarios, we decided to split the data into blocks of 100 frames and take the first 80 for training and the last 20 for testing. Therefore our model should have seen some city driving.

---

[5] Indeed, the number of parameters decreased from a total of 636.225 trainable parameters to 156.225 — a decrease by a factor of 4. We calculated these numbers using the tool *PyTorch summary* (https://github.com/sksq96/pytorch-summary).

[6] https://pytorch.org/docs/stable/index.html

VI. FURTHER IDEAS

A. Augmented brightness

The calculation of the optical flow field is in general quite sensitive to noise and brightness changes. To make these calculations more robust, we tried to add some additional noise to the frames, before calculating the flow field. To change the brightness and contrast of an image, we used the formula

$$\text{frame}_{\text{augmented}}(i,j) = \alpha(i,j) \cdot \text{frame}(i,j) + \beta(i,j).$$

The function α controls the contrast augmentation (> 1 increase, < 1 decrease) and the function β controls the brightness. To get some noise into the frames, we used

$$\alpha \sim \mathcal{U}(0,1) + 0.35, \ \beta \sim \mathcal{U}(-5,35),$$

where $\mathcal{U}(a,b)$ denotes the uniform distribution in an interval $[a,b]$ for $a,b \in \mathbb{R}$ with $a < b$.

B. Increase data

Use MUCH MUCH MUCH MUCH MUCH more data!
Create Validation Data that is not connected to training data and still covers all situations

REFERENCES

[1] Markus Maurer et al. *Autonomous Driving: Technical, Legal and Social Aspects*. Berlin, Heidelberg: Springer, 2016.
[2] Gunnar Farnebäck. "Two-Frame Motion Estimation Based on Polynomial Expansion". In: *Scandinavian Conference on Image Analysis* (2003), pp. 363–370.
[3] Mariusz Bojarski et al. "End to End Learning for Self-Driving Cars". In: (Apr. 2016). URL: https://arxiv.org/pdf/1604.07316v1.pdf.
[4] Sen Wang et al. "DeepVO: Towards end-to-end visual odometry with deep Recurrent Convolutional Neural Networks". In: *2017 IEEE International Conference on Robotics and Automation (ICRA)* (May 2017). DOI: 10.1109/icra.2017.7989236. URL: http://dx.doi.org/10.1109/ICRA.2017.7989236.
[5] Sergey Ioffe and Christian Szegedy. "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift". In: (Feb. 2015). URL: https://arxiv.org/pdf/1502.03167.pdf.
[6] Nitish Srivastava et al. "Dropout: A Simple Way to Prevent Neural Networks from Overfitting". In: *Journal of Machine Learning Research* 15.56 (2014), pp. 1929–1958. URL: http://jmlr.org/papers/v15/srivastava14a.html.
[7] Diederik Kingma and Jimmy Ba. "Adam: A Method for Stochastic Optimization". In: *International Conference on Learning Representations* (Dec. 2014).