

A machine learning approach to predict a vehicles velocity using dashcam video

Florian Wolf

Department of Mathematics and Statistics, University Konstanz
Konstanz, Germany
florian.2.wolf@uni-konstanz.de

Franz Herbst

Department of Physics, University Konstanz
Konstanz, Germany
franz.herbst@uni-konstanz.de

Abstract—Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

Index Terms—deep learning, computer vision, visual odometry, dense optical flow, siamese network

I. INTRODUCTION

Autonomous driving is considered to have a key role in the future of human mobility. Recently the topic has therefore gained great attention in research, economy and politics [Maurer2016]. In this project we approached this topic in a related but more simplified setup and tried out several machine learning techniques and network architectures.

A. Aim of the project

In this project we tried to predict the velocity of a car based on a dashcam video. This was inspired by the *comma*¹ speed challenge² published in 2018. To achieve this goal we used optical flow analysis to preprocess the data and neuronal networks with different structures (classical and siamese) to predict the velocity. In order to measure the quality of the predictions we used their mean squared errors of the measured velocity.

¹<https://comma.ai/>

²<https://github.com/commaai/speedchallenge>

II. DATA COLLECTION, ANALYSIS AND PREPROCESSING

To start the project we used the database given in the *comma* speed challenge project, which consists of a 17 minute training video (20400 frames) and the corresponding car velocity, as well as a 9 minute test video without labels to validate the model on an unknown data set.

A. Data collection

Due to our limited computational power we mainly used this data set to study different training techniques and model architectures. Still we could not expect good generalization for this rather small pool of training data. We therefore developed a technique to acquire more data that needed minimal resources. We used the open source smartphone apps *open camera* and *open street maps* to cast a video while driving and map the velocity using GPS at the same time. So we were able to train our final model with 1 h 40 min of video data.

B. Data analysis

In order to evaluate our raw data with the model we analysed our dataset concerning the recorded images and driving scenarios. The original frames have a size of $640 \times 380 \times 3$ pixels (RGB). To reduce the computation time, we cut the borders of the frames to exclude parts not needed for detection and sampled them down to half of their pixel size.

To visualize the training process in different driving scenarios, we plotted the velocity in dependence of the frame number (Fig. 1). After additional evaluation of the video we were able to classify three different driving scenarios: *highway driving* with high and relatively steady velocities, *stop and go* with low and fluctuating velocities as well as *city driving* with very abrupt speed changes between medium and very low velocities. Every driving scenario is represented with about one third of the training data.

We would then evaluate the predictions of each model $p(x_i)$ to the input data x_i using the mean squared error

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N (p(x_i) - y_i)^2 \quad (1)$$

to real value of the velocity y_i . To classify the resulting value we on one hand used the velocity plot and the following rough




Fig. 1: Distribution of the velocities of all frames in the training video, including the three categories.

classification³: $\sqrt{\mathcal{L}} \gtrsim 16$ *no fitting*, $16 \gtrsim \sqrt{\mathcal{L}} \gtrsim 10$ *average velocity fitted*, $10 \gtrsim \sqrt{\mathcal{L}} \gtrsim 5$ *scenarios detected*, $5 \gtrsim \sqrt{\mathcal{L}} \gtrsim 1$ *variances within scenarios detected*, $1 \gtrsim \sqrt{\mathcal{L}}$ *perfect fitting*.

For the training process we used a splitting of 80% training data and 20% validation data. Initially we did a hard splitting of the entire data, but this neglects the distribution of different driving scenarios in the video. Therefore we went on first splitting the data into driving scenarios with then assigning training data and validation data on each.

C. Preprocessing

After the previously explained preparation steps we needed to find a method to extract information about the motion of the vehicle to be able to predict its velocity. There are two ways of doing this: either fit two following frames into a network, which will be discussed later, or calculate what has changed between two frames and fit the results into a model.

To do this we used the “Farneback pyramid method” [Farneback2003] to calculate the optical flow.

with the following parameters

pyramid levels := 3
pyramid scaling := 0.5
window size := 6
pixel neighborhood size := 5
SD of the gaussian filter := 1.1

We choose three pyramid levels, because we wanted the calculations to be more accurate. To decrease the training duration, we halved the size of the optical flow frames again, resulting in a resolution of (160, 105, 3) pixels per frame. As we used a window size of six pixels, a comparison between the original optical flow and the down sampled one lead to the result, that we do not loose a lot of information.

³These rules do only apply on datasets with an equal distribution of all driving scenarios

The optical flow calculation returns the magnitude and the angle of the flow vectors, which we transformed into polar coordinates. To get an RGB image representing the optical flow of two consecutive frames, we normalized the magnitudes and put them into the third channel of the frame. The values of the second channel were all set to the value 255. We then multiplied the angle with the factor $180/(2\pi)$ and set this value for the first channel.

We wanted to see, if the model performs better using the dashcam frames as additional material. Therefore, we did the same down sampling with the frames.

III. METHOD SELECTION AND ARCHITECTURE

The prediction of the vehicles speed is a non-linear regression task, so the choice of a neural network is reasonable. Recent architectures (ResNet, GoogLeNet, etc.) have shown that using multiple stacked convolution layers combined with stacked dense layers, perform well on image classification tasks. Therefore the choice of a convolutional neural network is justified.

A. Initial Network

As an initial architecture we used the model presented by the NVIDIA work group [NVIDIA2016]. As it was designed for self-driving cars, the model has enough complexity to handle a task like ours and allowed various possibilities to fine-tune and improve the model. The raw structure is shown in Fig. 2.




Fig. 2: Initial Network using a series of convolutional layers which are then flattened and

Using the initial model on the training data with a hard splitting we achieved a MSE of 18-20 in the validation set while having a MSE of less than 3 on the training set. The initial network therefore has some clear overfitting problems, which we tried to address with our fine tuning.

B. Siamese Network

In order to improve our results we also tried to expanded our model and use the original image as well. The corresponding model is shown in Fig. 3.




Fig. 3: Siamese Network

We could use this architecture in two different ways: feeding two consecutive frames into the model or one image and one

optical flow field. In order to keep the structure correspondent we sampled down the regarding images to the size of the optical flow field. inspired by <https://arxiv.org/pdf/1709.08429.pdf>

IV. FINE TUNING

A. Batch Normalisation

To speed up the training we used batch normalization layers [BatchNorm2015]

B. Activation Function

and we tested different activation functions, to improve the performance of the model.

We used ReLu

$$\text{ReLU} : \mathbb{R} \rightarrow \mathbb{R}_0^+, x \mapsto \max\{0, x\}$$

as the initial activation function. Using the ReLU function and 15 epochs for training, we achieved a MSE of around 15 on the testing set. We ran the code multiple times, to ensure this result holds. This result was not really promising, so we wanted to decrease the error by modifying the model even more.

To solve the problem of dead neurons⁴ of the ReLU function, we tried the leakyReLU function

$$\text{leakyReLU} : \mathbb{R} \rightarrow \mathbb{R}, x \mapsto \begin{cases} x, & x \geq 0 \\ c \cdot x, & x < 0 \end{cases}$$

with a hyperparameter $c = 0.01$. Using the leakyReLU function, we achieved a MSE of around 12 on the testing set and under 3 on the training set. All of the models were trained with only eight epochs, as we still had problems with overfits.

C. Dropout Layer

As we still had a lot of issues with overfits, we decided to include a dropout layer, according to [Dropout2014]. We tried different numbers and positions for dropout layers, but using one layer with a dropout probability of $p = 0.5$ after the third convolutional layer seemed to work best.

D. Pooling layers with initial splitting

We added two generic pooling layers to the network to reduce the number of parameters of the model. One after the second convolutional layer and the second one right before the fully connected layers start⁵. We tested maximum and average pooling with the following parameters

$$\begin{aligned} \text{kernel size} &:= 2 \times 2 \\ \text{stride} &:= 2 \\ \text{padding} &:= 1 \\ \text{dilation} &:= \text{None} \end{aligned}$$

⁴One can clearly see in the definition of the ReLU function, that neurons with a value below zero cannot participate in the learning process.

⁵Indeed, the number of parameters decreased from a total of 636.225 trainable parameters to 156.225. Therefore the number of parameters decreased by a factor of 4. We calculated these numbers using the tool ‘PyTorch summary’(<https://github.com/sksq96/pytorch-summary>).

Initial splitting, 8 epochs	ReLU		leakyReLU	
	Train	Test	Train	Test
No pooling	2.85	12.08	2.45	10.75
Max pooling	5.62	11.82	5.52	10.29
Max pooling (15 epochs)	-	-	3.22	9.63
Average pooling	7.70	11.40	6.08	13.09

TABLE I: MSE results of the network using different pooling strategies, one dropout layers, two different activation functions and the initial splitting. We trained each of the models for eight epochs.

The implementation of pooling layers helped a lot, as now the loss on the train and test data seemed to decrease nearly equally. Our results with the initial splitting are shown in Table I. We gave the network with max pooling a try with 15 epochs, as the loss on the train and test set was decreasing in a pretty stable manner. Using this network, we achieved for the first time a MSE of under 10 on the test set, which is according to our assumptions in Section I-A a pretty good result — especially, as we trained the model mostly on highway scenes and tested it only in city driving scenarios.

V. TRAININGPROCESS

we used bla bla as optimizers and bla bla as schedulers, for each model we tried different combinations and bla bla

VI. ERROR ANALYSIS AND RESULTS

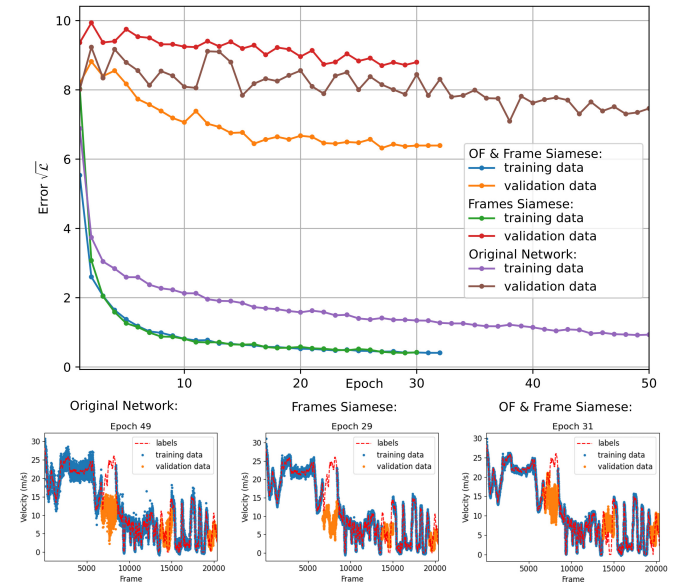


Fig. 4: Comparison of the results using the models: tuned original network and siamese model with frames and optical flow; *top*: training and validation error per epoch; *bottom*: predictions after the last epoch with correct values in red

We identified three possible reasons for our poor results:

- (i) predictions are limited to specific video parameters: predictions are bound to camera perspective, changes in brightness
- (ii) lack of generalization: perfect fitting on randomly shuffled validation data, qualitative fitting for block splitting and rough fitting for different videos
- (iii) very complex model trained with very little data

We came up with the following approaches to solve these problems

- (i) We wanted to try adding some additional noise into the frames before calculating the optical flow, to make the calculation more robust against brightness changes. As intentionally adding noise to a frame is quite atypical in computer vision, this idea looked quite interesting.
- (ii) Use another splitting. To get a better ratio between highway and city driving scenarios, we decided to split the data into blocks of 100 frames and take the first 80 for training and the last 20 for testing. Therefore our model should have seen some city driving.

VII. FURTHER WORK

A. *Augmented brightness*

STILL UNDER CONSTRUCTION

B. *Increase data*

Use MUCH MUCH MUCH MUCH MUCH more data!

Create Validation Data that is not connected to training data and still covers all situations

ACKNOWLEDGMENT

We like to thank bla bla bla