

Lecture 10 Clustering

Dr. Hanhe Lin

Dept. of Computer and Information Science
University of Konstanz




Unsupervised learning

- In unsupervised learning, the training data consists of a set of input vectors without any corresponding target values
- Goals:
 - Clustering: discover groups of similar examples within the data
 - Density estimation: determine the distribution of data within the input space
 - Data visualization: project the data from a high-dimensional space down to two or three dimensions

Unsupervised learning

- Unsupervised learning was largely ignored by machine learning community
 - It is hard to say what the aim of unsupervised learning is
- Unsupervised learning is the future of machine learning
 - Large scale of unlabeled data
 - Very expensive and infeasible to annotate large data

Clustering - intuition



Clustering

- Goal: partition the data into different groups (clusters), where the data in each group are similar to each other other than to those in other groups
- Clustering is one of the most commonly researched unsupervised learning topic
- The only information clustering uses is the similarity between examples
- Clustering groups examples based of their mutual similarities
- A good cluster algorithm is:
 - High within-cluster similarity
 - Low inter-cluster similarity

Clustering: when and why?

- Useful for:
 - Automatically organizing data
 - Understanding hidden structure in some data
 - Representing high-dimensional data in a low-dimensional space
- Examples:
 - Image segmentation
 - Customers according to purchase histories
 - Genes according to expression profile
 - Search results according to topic
 - Social network analysis

Clustering notation

Our data are

$$D = \{x^{(1)}, \dots, x^{(m)}\}$$

Each data point is n -dimensional, i.e.,

$$x = \{x_1, \dots, x_n\}$$

Define a distance function between data $x^{(i)}$ and $x^{(j)}$, $d(x^{(i)}, x^{(j)})$

Goal: segment the data into k groups

Clustering division

- Hierarchical clustering
 - Agglomerative: bottom-up
 - Divisive: top-down
 - ...
- Centroid-based clustering
 - K-means clustering
 - Kernel k-means
 - Fuzz c-means
 - ...
- Distribution-based clustering
 - Gaussian mixture model (GMM)
 - ...
- Density-based clustering
 - Mean-shift
 - ...

Distance measure


- Distance is inversely related to similarity, namely, large distance, low similarity; small distance, high similarity
- Choice of the distance measure is very important for clustering
- Some commonly used distance metrics:
 - Euclidean distance (L2 norm): $d(x^{(i)}, x^{(j)}) = \|x^{(i)} - x^{(j)}\|_2$
 - Manhattan distance (L1 norm): $d(x^{(i)}, x^{(j)}) = \|x^{(i)} - x^{(j)}\|_1$
 - Mahalanobis distance: $d(x^{(i)}, x^{(j)}) = \sqrt{(x^{(i)} - x^{(j)})^T S^{-1} (x^{(i)} - x^{(j)})}$
where S is the covariance matrix
- Similarity is subjective and hard to define
- Different similarity criteria can lead to different clusters

The property of a distance measure

- Non-negative: $d(x^{(i)}, x^{(j)}) \geq 0$
- Symmetry: $d(x^{(i)}, x^{(j)}) = d(x^{(j)}, x^{(i)})$
- Self-similarity: $d(x^{(i)}, x^{(j)}) = 0$ if $x^{(i)} = x^{(j)}$, we should not conclude A looks like B, but B does not look like A
- Triangle inequality: $d(x^{(i)}, x^{(k)}) + d(x^{(k)}, x^{(j)}) \geq d(x^{(i)}, x^{(j)})$


Agglomerative clustering

- Initially, each point is treated as a cluster
- Repeat:
 - Compute distance between all clusters in terms of the defined distance measure (store for efficiency)
 - Merge two nearest clusters into a new cluster
 - Stop when there is only one cluster left
- Save both clusters and sequence of cluster operations by “Dendrogram”




Agglomerative clustering


- How should we define “closest” for clustering with multiple elements?
- Many options:
 - **Closest pair** (single-link clustering)
 - **Farthest pair** (complete-link clustering)
 - Average of all pairs
- Different choices create different clustering behaviors



Agglomerative clustering





Closest pair



Farthest pair


Example 1: agglomerative clustering



Dendrogram

Example 2: agglomerative clustering

- Clustering gene expression data [Eisen et al 1998].



Overview: agglomerative clustering

- Summary
 - Choose a cluster distance/ dissimilarity method
 - Successively merge closet pair of clusters until only one cluster left
- Pros and cons
 - Easy to understand and implement
 - Don't need to define the number of clusters
 - Possible to view partitions at different levels of granularities
 - Computational extensive, complexity is $O(m^2 \log m)$


k-means clustering

- The basic idea is to describe each cluster by its mean value.
- Note: this works only for distances such that a mean is well-defined.
- The goal of k-means is to assign data to clusters and define these clusters with their means.


k-means clustering

- Input:
 - Number of clusters k
 - Data set $\{x^{(1)}, \dots, x^{(m)}\}$
- Steps:
 - Randomly initialize k cluster centroids, i.e., mean/average
 - Repeat until converge
 - Assign each data to its closest centroid
 - Change the cluster center to the average of its assigned points
- Output:
 - Centroids of each cluster
 - The cluster label which is assigned to each data


k-means clustering




k-means clustering




k-means clustering




k-means clustering




k-means clustering




k-means clustering




k-means clustering




k-means clustering



k-means clustering



k-means clustering



k-means clustering

Let be μ_1, \dots, μ_K the centriods of clusters

Optimization objective:

$$\min_{\mu_k, b_k^{(i)}} \sum_{i=1}^m \sum_{k=1}^K b_k^{(i)} \|x^{(i)} - \mu_k\|^2$$

$$\text{where } b_k^{(i)} = \begin{cases} 1 & \text{if } \|x^{(i)} - \mu_k\| = \min_j \|x^{(i)} - \mu_j\| \quad j = 1, \dots, K \\ 0 & \text{otherwise} \end{cases}$$

$b_k^{(i)}$ is the indicator to represent if the data i is assigned to cluster k

Exact optimization of the k-means objective is NP-hard

k-means clustering

The solution for k-means algorithm is heuristic

Assign step: fix centroids μ_1, \dots, μ_K , optimize $b_k^{(i)}$

Mean relocation step: fix $b_k^{(i)}$, optimize centroids μ_1, \dots, μ_K

Each step is guaranteed to decrease the objective, thus guaranteed to converge to a local minimum.

k-means is a coordinate descent algorithm.




k-means clustering

Algorithm:

- Randomly initialize k cluster centroids
- Repeat:
 - For 1 to m
 - Compute its indicator
 - For 1 to k
 - Compute the new cluster centroids

Random initialization

- k-means algorithm converges to a local minimum, so it is extremely sensitive to cluster center initialization.
- Bad initialization may lead to:
 - Poor convergence speed
 - Bad overall clustering





Handle bad initialization

- Make initialized centroids evenly distribute in training set:
 - Choose first center as one of the examples, second which is the farthest from the first, third which is the farthest from both, and so on
- Try-and-see, namely, try multiple initializations and choose the best result:
 - For $i = 1$ to 100
 - Run k-means
 - Compute cost function
 - Pick the result that gave the lowest loss

$$L = \sum_{i=1}^m \sum_{k=1}^K b_k^{(i)} \|x^{(i)} - \mu_k\|^2$$

Choosing the number of cluster k

- Elbow method: try different values of k , plot the k-means loss function L , and find out the “elbow-point” in the figure




Example: image segmentation by k-means



Figure 9.3 Two examples of the application of the K -means clustering algorithm to image segmentation showing the initial images together with their K -means segmentations obtained using various values of K . This also illustrates the use of vector quantization for data compression, in which smaller values of K give higher compression at the expense of poorer image quality.

k-means clustering: summary

- k-means is the most popular and widely used clustering algorithm
- Converge to a local minimum
- Computational complexity in each iteration:
 - Assign data points to closest cluster center $O(mk)$
 - Change the cluster center to the average of its assigned points $O(m)$
- “Hard assignment”, either 1 or 0 to each cluster
- Limitations:
 - Sensitive to outliers
 - Works well only for round shaped, and of roughly equal sizes/density clusters
- Do not work well on non-linear clusters



k-medoids

- In many practical settings, Euclidean distance is not appropriate
- For example:
 - Discrete multivariate data, such as purchase histories
 - Positive data, such as time spent on a web-page
- k-medoids is an algorithm that only requires knowing distances between data points $x^{(i)}$ and $x^{(j)}$, $d(x^{(i)}, x^{(j)})$
- No need to define the mean
- Each of the clusters is associated with its most typical example

k-medoids clustering

Algorithm:

- Randomly choose k data as initial cluster center
- Repeat:
 - Assign each data point to its closest center
 - For each cluster, find the data point in that cluster that is closest to the other points in that cluster
 - Set each cluster center equal to their closest data points

Agglomerative vs. k-means


Agglomerative:

- Gives different partitioning depending on the level-of-resolution we are looking at
- Doesn't need the number of clusters to be specified
- May be slow as it has to compute distance many times $O(m^2 \log m)$


k-means:

- Produces a single partitioning
- Needs the number of clusters to be specified
- More efficient $O(m(k + 1)i)$, where i is the number of iteration

GMM



GMM




Gaussian distribution

The probability density of the normal (Gaussian) distribution is

$$p(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

where μ is the mean, and σ^2 is the variance



Multivariate Gaussian distribution

The probability density function of multivariate (n -dim) Gaussian distribution is given by:

$$p(x|\mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1} (x-\mu)}$$



where μ is the mean, Σ is the covariance matrix, and $|\Sigma|$ denotes the determinant of the covariance matrix

Partial derivative to mean and covariance:

$$\frac{\partial}{\partial \mu} p(x|\mu, \Sigma) = p(x|\mu, \Sigma) \Sigma^{-1} (x - \mu)$$



$$\frac{\partial}{\partial \Sigma} p(x|\mu, \Sigma) = -\frac{1}{2} p(x|\mu, \Sigma) (\Sigma^{-1} - \Sigma^{-1} (x - \mu)(x - \mu)^T \Sigma^{-1})$$

Multivariate Gaussian distribution





$$\mu = \begin{bmatrix} 0 & 0 \end{bmatrix} \quad \Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

Multivariate Gaussian distribution



$$\mu = \begin{bmatrix} 0 & 0 \end{bmatrix} \quad \Sigma = \begin{bmatrix} 1 & 0.8 \\ 0.8 & 1 \end{bmatrix}$$

Multivariate Gaussian distribution



$$\mu = \begin{bmatrix} 3 & 3 \end{bmatrix} \quad \Sigma = \begin{bmatrix} 2 & -0.8 \\ -0.8 & 1 \end{bmatrix}$$

GMM – model representation

- Let $z^{(i)}$ be the latent variable, and $z^{(i)} \in \{1, 2, \dots, k\}$. $p(z^{(i)} = j)$ denotes the probability that $x^{(i)}$ generated from j th Gaussian model (k models in total)
- Gaussian mixture density is given by:

$$p(x^{(i)}) = \sum_{j=1}^k p(z^{(i)} = j) p(x^{(i)} | \mu^{(j)}, \Sigma^{(j)})$$

- Assume $z^{(i)} \sim \text{Multinomial}(\alpha)$, where $\alpha_j \geq 0$, $\sum_{j=1}^k \alpha_j = 1$. Correspondingly,
 $p(z^{(i)} = j) = \alpha_j$

- Gaussian mixture density is rewritten as:

$$p(x^{(i)}) = \sum_{j=1}^k \alpha_j p(x^{(i)} | \mu^{(j)}, \Sigma^{(j)})$$

where α_j is mixture coefficient

GMM

- The parameters of GMM are thus $\alpha_j, \mu^{(j)}, \Sigma^{(j)} \forall j = 1, 2, \dots, k$
- Given the training data $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$, Maximum Likelihood Estimation (MLE) maximizes the following log likelihood:

$$\begin{aligned} L &= \log \left(\prod_{i=1}^m p(x^{(i)}) \right) \\ &= \sum_{i=1}^m \log \left(p(x^{(i)}) \right) \\ &= \sum_{i=1}^m \log \left(\sum_{j=1}^k \alpha_j p(x^{(i)} | \mu^{(j)}, \Sigma^{(j)}) \right) \end{aligned}$$

GMM

If parameters $\alpha_j, \mu^{(j)}, \Sigma^{(j)} \forall j = 1, 2, \dots, k$ maximize log likelihood, then we have:

$$\frac{\partial L}{\partial \mu^{(j)}} = \sum_{i=1}^m \frac{\alpha_j}{\sum_{l=1}^k \alpha_l p(x^{(i)} | \mu^{(l)}, \Sigma^{(l)})} \frac{\partial}{\partial \mu^{(j)}} p(x^{(i)} | \mu^{(j)}, \Sigma^{(j)}) = 0$$

$$\frac{\partial L}{\partial \Sigma^{(j)}} = \sum_{i=1}^m \frac{\alpha_j}{\sum_{l=1}^k \alpha_l p(x^{(i)} | \mu^{(l)}, \Sigma^{(l)})} \frac{\partial}{\partial \Sigma^{(j)}} p(x^{(i)} | \mu^{(j)}, \Sigma^{(j)}) = 0$$

Combined with

$$\frac{\partial}{\partial \mu} p(x | \mu, \Sigma) = p(x | \mu, \Sigma) \Sigma^{-1} (x - \mu)$$

$$\frac{\partial}{\partial \Sigma} p(x | \mu, \Sigma) = -\frac{1}{2} p(x | \mu, \Sigma) (\Sigma^{-1} - \Sigma^{-1} (x - \mu)(x - \mu)^T \Sigma^{-1})$$

GMM

Bayes rule: $p(A | B)p(B) = p(B | A)p(A)$

Applied to GMM: $p(z^{(i)} = j | x^{(i)})p(x^{(i)}) = p(z^{(i)} = j)p(x^{(i)} | z^{(i)} = j)$

$$\Rightarrow p(z^{(i)} = j | x^{(i)}) \sum_{l=1}^k \alpha_l p(x^{(i)} | \mu^{(l)}, \Sigma^{(l)}) = \alpha_j p(x^{(i)} | \mu^{(j)}, \Sigma^{(j)})$$

Posterior probability:

$$\underbrace{p(z^{(i)} = j | x^{(i)})}_{\gamma_{ij}} = \frac{\alpha_j p(x^{(i)} | \mu^{(j)}, \Sigma^{(j)})}{\sum_{l=1}^k \alpha_l p(x^{(i)} | \mu^{(l)}, \Sigma^{(l)})}$$

The mean and variance are computed as:

$$\mu^{(j)} = \frac{\sum_{i=1}^m \gamma_{ij} x^{(i)}}{\sum_{i=1}^m \gamma_{ij}} \quad \Sigma^{(j)} = \frac{\sum_{i=1}^m \gamma_{ij} (x^{(i)} - \mu^{(j)})(x^{(i)} - \mu^{(j)})^T}{\sum_{i=1}^m \gamma_{ij}}$$


GMM

Bayes rule: $p(A | B)p(B) = p(B | A)p(A)$

Applied to GMM: $p(z^{(i)} = j | x^{(i)})p(x^{(i)}) = p(z^{(i)} = j)p(x^{(i)} | z^{(i)} = j)$

$$\Rightarrow p(z^{(i)} = j | x^{(i)}) \sum_{l=1}^k \alpha_l p(x^{(i)} | \mu^{(l)}, \Sigma^{(l)}) = \alpha_j p(x^{(i)} | \mu^{(j)}, \Sigma^{(j)})$$

Posterior probability:

 **Soft assignment**

$$\underbrace{p(z^{(i)} = j | x^{(i)})}_{\gamma_{ij}} = \frac{\alpha_j p(x^{(i)} | \mu^{(j)}, \Sigma^{(j)})}{\sum_{l=1}^k \alpha_l p(x^{(i)} | \mu^{(l)}, \Sigma^{(l)})}$$

The mean and variance are computed as:

$$\mu^{(j)} = \frac{\sum_{i=1}^m \gamma_{ij} x^{(i)}}{\sum_{i=1}^m \gamma_{ij}} \quad \Sigma^{(j)} = \frac{\sum_{i=1}^m \gamma_{ij} (x^{(i)} - \mu^{(j)})(x^{(i)} - \mu^{(j)})^T}{\sum_{i=1}^m \gamma_{ij}}$$

GMM

We still have an additional constraint: $\alpha_j \geq 0, \sum_{j=1}^k \alpha_j = 1$

Introduce the Lagrange multiplier:

$$J = L + \lambda(\sum_{j=1}^k \alpha_j - 1)$$

Take derivative, we have

$$\frac{\partial J}{\partial \alpha_j} = 0 \Rightarrow \alpha_j = \frac{1}{m} \sum_{i=1}^m \gamma_{ij}$$

GMM algorithm

Randomly initialize parameters $\alpha_j, \mu^{(j)}, \Sigma^{(j)} \quad \forall j = 1, 2, \dots, k$

Repeat until converge:

E-step. Compute posterior probability γ_{ij} :

$$\gamma_{ij} = \frac{\alpha_j p(x^{(i)} | \mu^{(j)}, \Sigma^{(j)})}{\sum_{l=1}^k \alpha_l p(x^{(i)} | \mu^{(l)}, \Sigma^{(l)})}$$

M-step. Update the parameters using the current γ_{ij} :

$$\alpha_j = \frac{1}{m} \sum_{i=1}^m \gamma_{ij} \quad \mu^{(j)} = \frac{\sum_{i=1}^m \gamma_{ij} x^{(i)}}{\sum_{i=1}^m \gamma_{ij}} \quad \Sigma^{(j)} = \frac{\sum_{i=1}^m \gamma_{ij} (x^{(i)} - \mu^{(j)})(x^{(i)} - \mu^{(j)})^T}{\sum_{i=1}^m \gamma_{ij}}$$

The General EM algorithm

Given a joint distribution $p(X, Z | \theta)$ the over observed variables X and the latent variables Z , governed by parameters θ , the goal is to maximize the likelihood function $p(X | \theta)$ with respect to θ .

- Choose an initial setting for the parameters θ
- Repeat until converge:
 - E-step: evaluate $p(X, Z | \theta)$ using current θ
 - M-step: evaluate θ given by Z

GMM vs. k-means

GMM (soft assignment)

- Randomly initialize parameters of k Gaussian distribution and mixture coefficients
- Repeat until converge:
 - E-step. Compute soft membership, i.e., posterior probability
 - M-step. Update the parameters using the current soft membership

k-means (hard assignment)

- Randomly initialize k cluster centroids
- Repeat until converge
 - Assign each data to its closest centroid
 - Update the cluster center using its current assigned points